

Oct 24, 2024 | 📅 RV Performance Event Sampling TG

Attendees: Beeman Strong Snehasish Kumar tech.meetings@riscv.org

Notes

- **Attendees:** Beeman, Bruce, Snehasish, Chun
- **Slides/video** [here](#)
- Recap freeze discussion:
 - Freeze on overflow preferred with precise attribution (Sspesa), freeze on interrupt preferred otherwise (using epc for attribution)
 - There are some other usages, like RR, that always prefer freeze on overflow
 - Counter freeze value:
 - Avoid counting events in handler, for kernel sampling (where priv mode filtering isn't sufficient)
 - Easy way to get precise attribution for arch events, but not uarch events
 - So moderate utility
- For counter freeze, propose that:
 - Can select whether freeze is on interrupt or overflow
 - Can select which counters freeze on a per-counter basis. Select a group of counters where all freeze if any in the group get interrupt/overflow
 - So one group of counters involved in freeze
- Counter freeze is not on by default, only enabled when extension is implemented and user opts in by selecting some counters to freeze
- CTR freeze
 - Same considerations apply (prefer overflow freeze with precise attribution, otherwise interrupt freeze)
 - LCOFIFRZ is defined to freeze on interrupt (since no precise attribution today)
 - Propose adding OFFRZ bit to support freeze on overflow
- Sspesa access
 - shpms* accessible to S-mode when counter delegation enabled, and then only delegated counters update them
 - So M-mode can use Sspesa when CDE=0
 - VS-mode requires SBI to access, just like counters
- Does counter/CTR freeze happen for non-delegated counters, when CDE=1?
 - Otherwise differs from Sspesa
 - *Will discuss next time*
- Do multiple overflows really happen?
 - We don't know of any real usages, save for multi-user (e.g., hypervisor and guest both profiling in guest)
 - Opted not to add HW (more shpms* CSRs) to support a rare occurrence (multiple overflows in small window) in a rare usage (multiple sampling users)
- Output to memory for Sspesa

- Doable, but would require more ISA (counter reload mechanism, OF clearing, etc). Propose not adding that now.
- Could allow reload only on subset of counters
 - Agreed. ISA would probably specify reload for all, but an implementation could support it only for a subset.
 - We haven't discussed whether Ssplcofi or Sspesa require precise sampling/attribution for all Zihpm counters
 - Each counter adds incremental cost. E.g., bit(s) in the ROB
 - Lean towards letting implementations make value judgment on added cost vs the consistency that SW prefers
 - Haven't we already added HW backoffs with Ssplcofi vs Sspesa
 - Yes, but those are different HW tradeoffs
- If known usages have only 1 counter sampling, maybe we only need precise sampling/attribution support for 1 counter?
 - *Interesting question, let's discuss more next time*
- SiFive has PC sampling in trace. Could leverage shpms* HW?
- **Out of time**, continue this discussion next time

Action items

